

# Building a Robot Judge: Data Science for the Law

## 6. Regression, Classification, and Regularization

Elliott Ash

# OLS Regression Baseline

- ▶ Consider the linear model

$$Y_i = X_i' \theta + \epsilon_i$$

where  $Y_i$  and all elements of  $X_i$  have been de-meanned and standardized to  $s.d. = 1$ .

- ▶ OLS assumptions:
  - ▶  $X_i$  uncorrelated with  $\epsilon_i$ 
    - ▶ Let's just assume this for now; will come back later.
  - ▶ Columns of  $X_i$  are not highly collinear.
    - ▶ In the case of word/n-gram frequency data, this is not a good assumption.

# Univariate OLS to Rank Predictive Features

- ▶ Consider the univariate regression

$$Y_i = \theta_w x_i^w + \epsilon_i$$

for each text feature  $w$  (e.g., relative word or n-gram frequency).

- ▶ Can be estimated with OLS.

- ▶ Can add fixed effects:

$$Y_i = \alpha_i + \theta_w x_i^w + \epsilon_i$$

where  $\alpha_i$  is a vector of dummy variables (fixed effects) for categories, e.g. location, year.

- ▶ Even better, can residualize  $Y$  and  $X$  on fixed effects before running any regressions.
  - ▶ That is, regress  $Y_i = \alpha_i + \epsilon_i$  and  $x_i^w = \alpha_i + \epsilon_i, \forall w$ , take residuals  $\tilde{Y}_i = Y_i - \hat{\alpha}_i$  and  $\tilde{x}_i^w = x_i^w - \hat{\alpha}_i$
  - ▶ then regress

$$\tilde{Y}_i = \theta_w \tilde{x}_i^w + \epsilon_i$$

# Estimating OLS

- ▶ OLS minimizes the mean squared error:

$$\min_{\hat{\theta}} \frac{1}{m} \sum_{i=1}^m (\mathbf{x}_i' \hat{\theta} - y_i)^2$$


where  $\mathbf{x}_i = (x_{i1}, x_{i2}, \dots, x_{ip})$

- ▶ This has a closed form solution

$$\hat{\theta} = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{Y}$$

- ▶ Most machine learning models do not have a closed form solution.
  - ▶ objective must be minimized using some other optimization algorithm.

# OLS Example: Refugee Resettlement and Employment Rates



OLS coefficient plots from Bansak et al, 2018 – effect of refugee characteristics on employment rates, separately by location.

# Gradient Descent





Figure 4-3. Gradient Descent


- ▶ Gradient descent measures the local gradient of the error function, and then steps in that direction.
  - ▶ Once the gradient equals zero, you have reached a minimum.

# Gradient Descent Pitfalls



- Gradient descent will find a local minimum, not necessarily a global minimum. And it can get stuck on plateaus.

# Gradient Descent and Scaling



- When using gradient descent, all features should be standardized to the same scale to speed up convergence.



# Gradient Descent Implementation

- ▶ The partial derivative of MSE for feature  $j$  is

$$\frac{\partial \text{MSE}}{\partial \theta_j} = \frac{2}{m} \sum_{i=1}^n (\mathbf{x}'_i \hat{\theta} - y_i) x_{ij}$$

- ▶ The gradient is the vector of these partial derivatives is


$$\nabla_{\theta} \text{MSE} = \begin{bmatrix} \frac{\partial \text{MSE}}{\partial \theta_0} \\ \frac{\partial \text{MSE}}{\partial \theta_1} \\ \vdots \\ \frac{\partial \text{MSE}}{\partial \theta_j} \end{bmatrix} = \frac{2}{m} \mathbf{X}' (\mathbf{X}' \theta - \mathbf{y})$$

- ▶ Gradient descent step:

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} \text{MSE}$$

# Stochastic Gradient Descent


- ▶ SGD picks a random instance in the training set and computes the gradient only for that single instance.



- ▶ Much faster to train, but can still bounce around even after it is close to the minimum.

# Mini-Batch Gradient Descent

- ▶ A compromise between gradient descent and stochastic gradient descent is mini-batch gradient descent, which selects a small number of rows (a “mini-batch”) for gradient compute, rather than a single row.



# Document Classification

- ▶ Say you have a set of cases where the defendant is held innocent or guilty.
  - ▶ You only know the outcome for a subset of documents which have been hand-coded.
  - ▶ Can we use the existing labels, and the text of all cases, to machine-code the labels for the unlabelled cases?
- ▶ sklearn has a number of classifiers for this purpose, e.g. LogisticRegression.

# Confusion Matrix

- ▶ Use **cross\_val\_predict()** to obtain a “clean” prediction for each row, in the sense that it is trained on data outside that row’s fold.
- ▶ A confusion matrix is a nice way to visualize classifier performance:

		Predicted Class	
		Negative	Positive
True Class	Negative	# True Negatives	# False Positives
	Positive	# False Negatives	# True Positives

- ▶ The values in the table give counts in the evaluation set.

# Precision and Recall

		Predicted Class	
		Negative	Positive
True Class	Negative	# True Negatives	# False Positives
	Positive	# False Negatives	# True Positives

- ▶ Two alternative metrics used to understand classifiers:

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

$$\text{Recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

- ▶ **Precision decreases with false positives.**
- ▶ **Recall decreases with false negatives.**

# F1 Score

- ▶ The  $F_1$  score provides a single combined metric – it is the harmonic mean of precision and recall:

$$F_1 = \frac{2}{\frac{1}{\text{precision}} + \frac{1}{\text{recall}}} = 2 \times \frac{\text{precision} \times \text{recall}}{\text{precision} + \text{recall}} =$$
$$= \frac{\text{Total Positives}}{\text{Total Positives} + \frac{1}{2}(\text{False Negatives} + \text{False Positives})}$$



- ▶ What about True Negatives?

# The Precision/Recall Tradeoff

- ▶ In general, one can tweak a classifier to increase precision at the cost of reducing recall, and vice versa.
  - ▶ The F1 score values them symmetrically.
- ▶ One can imagine contexts where they should be valued asymmetrically:
  - ▶ in the case of deciding “guilty” in court, you might prefer a model that lets many actual-guilty go free (high false negatives  $\leftrightarrow$  low recall) but has very few actual-innocent put in jail (low false positives  $\leftrightarrow$  high precision).
  - ▶ in the case of detecting bombs during flight screening, you might prefer a model that has many false alarms (low precision) to minimize the number of misses (high recall).




# The Precision/Recall Tradeoff



# ROC Curve and AUC

- Plots true positive rate (recall) against the false positive rate ( $FP / (FP + TN)$ ):



## Application: Facts vs. Law

- ▶ A crucial distinction in law practice, and in applying caselaw, is to distinguish the law from the facts.
- ▶ In Cao, Ash, Chen (2018), we made a sample of fact and law sections in court cases and trained a model to predict them based on the text.
- ▶ Corpus:
  - ▶ all cases from courtlistener.com with an annotated “FACTS” section (extracted using regular expressions).
  - ▶ 23,497 sections, split into 1.3M paragraphs
  - ▶ 36% facts, 64% laws

# Facts vs. Law: Supervised Learning Methods

	F1 for facts	F1 for values
DEPN	73.38	74.66
DEPW	72.77	73.79
SMITH	71.85	71.4
WS	<b>77.11</b>	<b>77.67</b>
BOW	72.57	73.29
D2V	67.18	65.36

Table 2: 5-fold cross validation.

**DEPN:** Document represented as a sequence of word-dependency pairs, fed to a MLP classifier with two 500-dim hidden layers.

**DEPW:** a variant of DEPN using Shulayeva et al.'s (2017) dependency features.

**SMITH** : Implementation of Smith's (2014) method: Bag of Words, but filter words that are statistically related to either fact statements or law statements. Fed to a Logistic Regression classifier.

**WS:** Laver et al.'s (2003) Wordscore algorithm (cf. Sarel and Demirtas, 2017). Bag of Words, but each word assigned a score based on relation to fact or law. Document is represented as a (re-scaled) score that sums up the pre- calculated scores of the words it contains, weighted by their frequencies.

**BOW** : Bag of words representation of documents, fed to a Logistic Regression classifier.

**D2V:** 500-dim Doc2Vec vectors (Le and Mikolov, 2014), fed to MLP.

# The Problem of Overfitting




Figure 4-14. High-degree Polynomial Regression

# Bias-Variance Tradeoff

- ▶ *Bias*

- ▶ Error due to wrong assumptions, such as assuming data is linear when it is quadratic.
- ▶ Will underfit the training data

- ▶ *Variance*

- ▶ Error due to excess sensitivity to small variations in the training data.
- ▶ A model with high variance is likely to overfit the training data.

- ▶ *Irreducible error*

- ▶ Error due to noise in the data.

# Training Notes

- ▶ In general, increasing a model's complexity will increase variance and reduce bias.
- ▶ If the model is underfitting:
  - ▶ adding more training data will not help.
  - ▶ use a more complex model
- ▶ If the model is overfitting:
  - ▶ adding more training data may help
  - ▶ or use regularization

# Lasso, Ridge, and Elastic Net

- ▶ Lasso and ridge regression are tools for dealing with large feature sets where:
  - ▶ models have multicollinearity that causes bias
  - ▶ models tend to overfit
  - ▶ models are computationally costly to fit
- ▶ These algorithms work by constraining estimated parameter sizes.



# Lasso Regression

- ▶ The Lasso cost function is

$$J(\theta) = \text{MSE}(\theta) + \alpha_1 \sum_{i=1}^n |\theta_i|$$

- ▶  $i$  indexes over  $n$  features
  - ▶  $\alpha_1$  is a hyperparameter setting the strength of the L1 penalty
- ▶ Lasso automatically performs feature selection and outputs a sparse model.

# Ridge Regression

- ▶ The Ridge cost function is

$$J(\theta) = \text{MSE}(\theta) + \alpha_2 \frac{1}{2} \sum_{i=1}^n \theta_i^2$$

- ▶  $i$  indexes over  $n$  features
  - ▶  $\alpha_2$  is a hyperparameter setting the strength of the L2 penalty
- ▶ It turns out that the Ridge estimator, like OLS, has a closed-form solution:

$$\hat{\theta}_{\text{Ridge}} = (X'X + \alpha_2 \mathbf{I}_n)^{-1} X' \mathbf{y}$$

where  $\mathbf{I}_n$  is the identity matrix.

- ▶ But it can also be solved by (stochastic) gradient descent.

# Elastic Net

- ▶ Elastic Net uses both L1 and L2 penalties:

$$J(\theta) = \text{MSE}(\theta) + \alpha_1 \sum_{i=1}^n |\theta_i| + \alpha_2 \frac{1}{2} \sum_{i=1}^n \theta_i^2$$

- ▶ in general, elastic net is preferred to lasso, which can behave erratically when the number of features is greater than the number of rows, or when some features are highly collinear.
- ▶ For elastic net, as with Lasso and Ridge, the hyperparameters can be selected by cross-validation.

## Scaling while maintaining sparsity

- ▶ Regularization penalties are designed to work with scaled data.
  - ▶ An important feature of text data is sparsity, which is lost when taking out the mean:

$$\tilde{x}_i = \frac{x_i - \bar{\mathbf{x}}}{\text{SD}[\mathbf{x}]}$$

- ▶ Solution:

$$\tilde{x}_i = \frac{x_i}{\text{SD}[\mathbf{x}]}$$


# Logistic Regression is for Classification

- ▶ Like OLS, logistic “regression” computes a weighted sum of the input features to predict the output.
  - ▶ But rather than output the sum directly, it transforms the sum using the logistic function.

$$\hat{p} = \Pr(Y_i = 1) = \sigma(\theta' \mathbf{x})$$

where  $\sigma(\cdot)$  is the sigmoid function

$$\sigma(a) = \frac{1}{1 + \exp(-a)}$$



- ▶ For the binary version of the classifier, the prediction  $\hat{Y} \in \{0, 1\}$  is determined by whether  $\hat{p} \geq .5$ .

# Logistic Regression Cost Function

- ▶ The cost function to minimize is

$$J(\theta) = \underbrace{-\frac{1}{m}}_{\text{negative}} \sum_{i=1}^m \left[ \underbrace{y_i}_{y_i=1} \underbrace{\log(\hat{p}_i)}_{\log \text{ prob}_{y_i=1}} + \underbrace{(1-y_i)}_{y_i=0} \underbrace{\log(1-\hat{p}_i)}_{\log \text{ prob}_{y_i=0}} \right]$$

- ▶ this does not have a closed form solution
- ▶ but it is convex, so gradient descent will find the global minimum.
- ▶ Just like linear models, logistic can be regulated with L1 or L2 penalties, e.g.:

$$J_2(\theta) = J(\theta) + \alpha_2 \frac{1}{2} \sum_{i=1}^n \theta_i^2$$

# Multi-Class Models

- ▶ Many interesting machine learning problems involve multiple un-ordered categories:
  - ▶ categorizing a case by area of law
  - ▶ predicting the political party of a speaker in a proportional representation system

# Multi-Class Confusion Matrix

		Predicted Class		
		Class A	Class B	Class C
True Class	Class A	Correct A	A, classed as B	A, classed as C
	Class B	B, classed as A	Correct B	B, classed as C
	Class C	C, classed as A	C, classed as B	Correct C

- More generally, can have a confusion matrix  $M$  with items  $M_{ij}$  (row  $i$ , column  $j$ ).



## Multi-Class Performance Metrics

Confusion matrix  $M$  with items  $M_{ij}$  (row  $i$ , column  $j$ ).

$$\text{Precision for } k = \frac{\text{True Positives for } k}{\text{True Positives for } k + \text{False Positives for } k} = \frac{M_{kk}}{\sum_j M_{kj}}$$

$$\text{Recall for } k = \frac{\text{True Positives for } k}{\text{True Positives for } k + \text{False Negatives for } k} = \frac{M_{kk}}{\sum_i M_{ik}}$$

$$F_1(k) = 2 \times \frac{\text{precision}(k) \times \text{recall}(k)}{\text{precision}(k) + \text{recall}(k)}$$

# Metrics for whole model

- ▶ Macro-averaging:
  - ▶ average of the per-class precision, recall, and F1, e.g.

$$F_1 = \frac{1}{n} \sum_{k=1}^n F_1(k)$$

- ▶ treats all classes equally
- ▶ Micro-averaging:
  - ▶ Compute model-level sums for true positives, false positives, and false negatives; compute precision/recall from model sums.

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}, \text{Recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

$$F_1 = 2 \times \frac{\text{precision} \times \text{recall}}{\text{precision} + \text{recall}}$$

- ▶ favors bigger classes
- ▶ “Weighted”: same as macro averaging, but classes are weighted by number of true instances in data.

# Multinomial Logistic Regression

- ▶ Logistic can be generalized to multiple classes.
  - ▶ When given an instance  $\mathbf{x}_i$ , multinomial logistic computes a score  $s_k(\mathbf{x}_i)$  for each class  $k$ ,

$$s_k(\mathbf{x}_i) = \theta'_k \mathbf{x}_i$$

- ▶ If there are  $n$  features and  $K$  output classes, there is a  $K \times n$  parameter matrix  $\Theta$ , where the parameters for each class are stored as rows.
- ▶ Using the scores, probabilities for each class are computed using the softmax function

$$\hat{p}_k(\mathbf{x}_i) = \Pr(Y_i = k) = \frac{\exp(s_k(\mathbf{x}_i))}{\sum_{j=1}^K \exp(s_j(\mathbf{x}_i))} = \frac{e^{\theta_k \mathbf{x}_i}}{\sum_{j=1}^K e^{\theta_j \mathbf{x}_i}}$$

- ▶ And the prediction  $Y_i \in \{1, \dots, K\}$  is determined by the highest-probability category.

# Multinomial Logistic Cost Function

- ▶ The binary cost function generalizes to the cross entropy

$$J(\theta) = \underbrace{-\frac{1}{m}}_{\text{negative}} \sum_{i=1}^m \sum_{k=1}^K \underbrace{\mathbf{1}[y_i = k]}_{y_i=k} \underbrace{\log(\hat{p}_k(\mathbf{x}_i))}_{\log \text{ prob}_{y_i=k}}$$

- ▶ again, this is convex, so gradient descent will find the global minimum.

## Application: Ash, Morelli, Osnabrugge 2018

- ▶ Use multinomial logistic regression to classify policy topics from plain text.

# Comparative Manifesto Project Corpus

- ▶ 44,020 annotated English-language political statements
  - ▶ hundreds of political party platforms from English-speaking countries.
- ▶ Each statement gets a CMP code, e.g. “decentralization”, “education”
  - ▶ 45 topics
    - ▶ some topics are somewhat esoteric, such as “marxist analysis”
  - ▶ We normalized to 19 broader, more interpretable topics

# Featurizing the Statements

- ▶ Each row includes a CMP code and a statement.
  - ▶ The statements are plain text.
- ▶ Standard featurization steps:
  - ▶ remove capitalization, punctuation, stopwords
  - ▶ construct n-grams up to length 3
  - ▶ remove n-grams appearing in less than 10 statements or more than 40 percent of statements
    - ▶  $M = 7,646$  features
  - ▶ compute tf-idf-weighted n-gram frequencies

# Regularized Logit Model

- ▶  $N$  rows,  $M$  text features,  $K$  policy topics
- ▶ Probability model:

$$P(Y_i = c) = \frac{e^{\beta_c X_i}}{\sum_{k=1}^K e^{\beta_k X_i}},$$

where  $c \in 1, \dots, K$  are the topic labels and  $\beta$  is an  $M \times K$  matrix of parameters.

- ▶ Cost function:

$$J(\beta) = -\frac{1}{N} \left[ \sum_{i=1}^N \sum_{k=1}^K \mathbf{1}\{y^i = k\} \log \frac{e^{\beta_k X_i}}{\sum_{l=1}^K e^{\beta_l X_i}} \right] + \gamma \sum_{j=1}^M \sum_{k=1}^K \beta_{jk}^2$$

- ▶  $\gamma$  = strength of L2 penalty
  - ▶  $\gamma^* = 1/2$  selected by 3-fold cross-validation grid search.



# Prediction Model

- ▶ Given a chunk of text, the logistic model computes a probability distribution over policy topics.
  - ▶ harnesses expert knowledge about political topics from Manifesto Project
- ▶ Validation of accuracy: predict the CMP code in a held-out sample of manifesto corpus statements
  - ▶ In-sample accuracy = 71%, Out-of-sample accuracy = 53%
    - ▶ quite good given there are 19 policy areas – choosing randomly would be correct 5% of the time; choosing top category (other topic) would be correct 15% of the time.

# Confusion Matrix

	Admin	Agric	Culture	Econ	Educ	Free- dom	Intl	Law / Order	Way of Life	Other	Party Politics	Quality of Life	Target Groups	Tech	Welfare	Total True
Administration	348	5	8	117	12	47	9	20	10	94	21	43	1	33	118	886
Agriculture	6	110	2	34	2	2	5	1	6	23	4	49	0	14	18	276
Culture	12	0	155	14	13	15	5	1	17	52	2	16	8	15	44	369
Economics	59	12	3	961	11	18	21	7	21	107	35	95	1	43	176	1570
Education	24	1	12	30	481	6	6	5	9	66	14	6	0	35	88	783
Freedom	51	0	10	24	0	240	29	14	23	69	28	3	6	2	59	558
Internationalism	15	3	2	41	2	16	453	12	17	75	14	21	4	8	44	727
Law and Order	18	1	0	19	10	13	24	361	9	83	12	2	1	14	58	625
Nat'l Way of Life	18	1	8	31	14	25	20	14	133	73	26	26	3	3	91	486
Other	55	18	15	128	61	34	52	64	35	1239	33	86	11	57	185	2073
Party Politics	19	3	3	51	10	18	17	7	22	90	181	25	1	5	59	511
Quality of Life	40	24	4	130	4	7	24	2	10	96	11	619	2	43	45	1061
Target Groups	16	1	7	13	13	7	8	10	13	35	3	1	57	7	71	262
Technology	28	7	6	63	29	3	6	4	8	73	7	52	0	397	48	731
Welfare	67	4	10	151	52	29	16	26	38	161	20	22	12	34	1454	2096
Total Predicted	776	190	245	1807	714	480	695	548	371	2336	411	1066	107	710	2558	13014

# Interpreting the Topics

- ▶ In the analysis of New Zealand parliamentary speeches, we looked at the n-grams that were predictive of the topic classification.
- ▶ Formally, for each topic  $k$  and each phrase  $m$ , regress

$$\Pr(y_i = k) = \alpha + \delta_m x_i^m + \epsilon_i$$

where  $\Pr(y_i = k)$  is the predicted probability that speech  $i$  is topic  $k$ , and  $x_i^m$  is the tf-idf frequency of phrase  $m$  in speech  $i$ .


- ▶ select phrases with highest positive t-statistic for  $\delta_m$

## Agriculture and Education Topics



# Dimensionality Reduction

- ▶ Especially in the case of text data, machine learning problems often involve thousands of features.
- ▶ Dimension reduction methods are needed:
  - ▶ not just for computational tractability, but also to help find a good solution.
  - ▶ also for data visualization – for example, to plot data in two dimensions.




The dimension reduction process matters: projecting down to two dimensions directly (left panel) might not isolate the variation we are interested in (as done in the right panel, which unrolls the Swiss Roll)

# Manifold Learning

- ▶ The swiss roll is an example of a 2D manifold:
  - ▶ like many real-world data sets, the data are not uniformly distributed across the space.
  - ▶ can be modeled in a lower-dimensional subspace while retaining most of the information
- ▶ Dimension reduction methods in machine learning are motivated by this “manifold hypothesis.”


# PCA



- ▶ PCA (principal components analysis), a popular dimension reduction technique.
  - ▶ identifies the axis that accounts for the largest amount of variance in the training set.
  - ▶ finds a second axis, orthogonal to the first, that accounts for the largest amount of the remaining variance, and so on
- ▶ The unit vector defining the  $i$ th axis is called the  $i$ th principal component.



# Choosing the number of dimensions




# Incremental PCA and numpy memmap

- ▶ Standard PCA requires you to load the whole data set into memory.
- ▶ numpy memmap loads big arrays from disk as needed
- ▶ sklearn.IncrementalPCA splits the data into mini-batches and trains gradually.

# Pros and Cons of PCA

- ▶ Advantages:
  - ▶ fast to compute
  - ▶ good performance on many tasks in practice
  - ▶ components are orthogonal by construction
- ▶ Disadvantages:
  - ▶ lose (potentially a lot of) predictive information from  $X$
  - ▶ Coefficients are not easily interpretable.
- ▶ Compromise:
  - ▶ keep strong predictors in feature set; take principal components of weak predictors

# t-SNE and MDS



From: L. Van der Maaten & G. Hinton, Visualizing Data using t-SNE, Journal of Machine Learning Research 9 (2008) 2579-2605


- ▶ **t-Distributed Stochastic Neighbor Embedding (t-SNE)** reduces dimensionality while trying to keep similar instances close and dissimilar instances apart.
  - ▶ Useful for visualizing clusters of instances in high-dimensional space
- ▶ **Multidimensional Scaling (MDS)** reduces dimensionality while trying to preserve the distances between the instances.

# Partial Least Squares


- ▶ Partial Least Squares is another technique to determine linear combinations of the predictive variables.
- ▶ Unlike PCA, the PLS technique works by successively extracting factors from both predictive and target variables such that covariance between the extracted factors is maximized.

# Ash 2016: PLS predictions of tax revenue changes using tax code text

## Income Tax




## Sales Tax



Weak predictors filtered out; 80% training, 20% testing sample.

- ▶ Predicted change in revenue (vertical axis), plotted against true change in revenue (horizontal axis).
- ▶ Correlations between truth and prediction: 0.89 and 0.84.



- ▶ The Google text classification guide recommends a maximum 20,000 items in your feature set.

## Feature selection using L1 Penalty

- ▶ A popular way to reduce dimensionality is to run lasso or elastic net, and exclude any predictors whose weights are regularized to zero.




# Feature selection using univariate comparisons

- ▶  $\chi^2$  is a very fast feature selection routine for classification tasks
  - ▶ features must be non-negative
  - ▶ works on sparse matrices
- ▶ With negative predictors, use `f_classif`.
- ▶ For regression tasks, use `f_regression`.
- ▶ There are also `mutual_info_regression` and `mutual_info_classif`, which will recover non-linear relations between predictor and outcome
  - ▶ they are much slower, can run on a sample of the data.

# Graph-based feature selection

- ▶ Graph-based feature selection works by clustering collinear features and then iteratively removing features.
  - ▶ much slower than univariate comparisons, but probably gives better performance in text classification tasks.
  - ▶ See Zhang and Hancock 2011.
  - ▶ There is an alpha implementation in python (not maintained) available at <https://github.com/danilkolikov/fsfc>.


# Voting Classifier



- ▶ voting classifiers generally out-perform the best classifier in the ensemble.
  - ▶ more diverse algorithms will make different types of errors, and improve your ensemble's robustness.


# Bagging and Pasting

- ▶ Rather than use the same data on different classifiers, one can use different subsets of the data on the same classifier:



- ▶ This is called **bagging** (bootstrap aggregating, when sampling with replacement) or **pasting** (when sampling without replacement).
- ▶ can also use different subsets of features across subclassifiers.
- ▶ The ensemble predicts by aggregating the predictions:
  - ▶ for classification, use the most frequent prediction
  - ▶ for regression, use the average output

# Bagging Benefits



- ▶ While the individual predictors have a higher bias than a predictor trained on all the data, aggregation reduces both bias and variance.
  - ▶ Generally, the ensemble has a similar bias but lower variance than a single predictor trained on all the data.
- ▶ Predictors can be trained in parallel using separate CPU cores.

# Random Forests

- ▶ Now you know how random forests work:
  - ▶ Random Forests are optimized ensembles of decision trees with bagging.
- ▶ Good prediction performance – due to out-of-sample validation being baked into the training process.


# Gradient Boosted Machines and XGBoost

- ▶ Gradient boosting works by sequentially adding predictors to an ensemble – it fits the new predictor to the residual errors made by the previous predictor to gradually improve the model.
- ▶ As of 2014, gradient boosted machines are considered better than random forests – they represent the state-of-the-art besides deep neural nets.

# Feature Importance

```
from xgboost import plot_importance  
plot_importance(xgb_reg, max_num_features=20)
```

<IPython.core.display.Javascript object>



1

- ▶ Random forests and boosted trees provide a metric of feature importance that summarizes how well each feature contributes to predictive accuracy.



## Ensemble Application: Jelveh, Kogut, and Naidu (2016)

- ▶ This paper looks at political language and ideology in the economics literature.
- ▶ They use data on campaign contributions to assign a subset of economists to Republican or Democrat.
- ▶ Then they train a classifier to predict party based on the text of written articles
  - ▶ They use an ensemble PLS model, that “votes” in the same way as random forests, but the constituent voters are PLS regressors, rather than decision trees.
  - ▶ They control for topic choices using JEL K codes and LDA topics.
  - ▶ The model predicts with 70% accuracy.

# JKN 2016: Results

- ▶ There is significant ideological sorting across fields:
  - ▶ law and economics is the most-right wing field, labor economics is the most left-wing field
- ▶ Right-wing economists report a higher labor supply elasticity than left-wing economists
- ▶ The ideology of editors does not affect ideology of published articles.