

Building a Robot Judge: Data Science for the Law

8. Deep Learning Essentials

Elliott Ash

“Neural Networks”

- ▶ “Neural”:
 - ▶ nothing like brains
- ▶ “Networks”:
 - ▶ nothing to do with “networks” as normally understood – in particular, nothing to do with network theory in math or social science.



Recent History

- ▶ NNs frequently outperform other ML techniques on very large and complex problems.
- ▶ Increase in computing power makes them computationally tractable, graphical processing units (GPUs, designed for video games) give you over 100x performance gain over CPUs.
- ▶ Training algorithms have improved – small tweaks have made a huge impact.
- ▶ Some theoretical limitations of ANNs have turned out to be benign in practice – for example, they work well on non-convex functions.


Will it last?

- ▶ Three key principles of deep learning will persist:
 - ▶ **Simplicity**
 - ▶ feature engineering is obsolete
 - ▶ complex, brittle, engineering-heavy pipelines replaced with simple, end-to-end trainable models, composed of 5-6 tensor operations.
 - ▶ **Scalability**
 - ▶ amenable to parallelization on GPUs or TPUs (tensor processing units)
 - ▶ trained on batches of data, so can be scaled to datasets of arbitrary size.
 - ▶ **Versatility and reusability**
 - ▶ can be trained on additional data without restarting from scratch, therefore amenable for continuous online learning.
 - ▶ deep-learning models are repurposable and thus reusable

Biological Neurons




Neurons can perform logical computations



- These networks perform the identity, logical AND, logical OR, and logical NOT operations.

Perceptron LTU



- In a perceptron, an individual neuron (called an LTU, or linear threshold unit) is defined by

$$h(\mathbf{x}) = \text{step}(\omega' \mathbf{x})$$


where

$$\text{step}(a) = \begin{cases} 0 & a < 0 \\ 1 & a \geq 0 \end{cases}$$

- The neuron computes a linear combination of the inputs; if result exceeds threshold, output positive class, otherwise negative class.

Perceptron

- ▶ A perceptron is an array of LTUs in parallel:



- ▶ This basic perceptron is similar to a logistic regression model.


In Notation

- ▶ The simplest perceptron is a linear combination of the inputs, the same as a linear regression:

$$y = \alpha + x'\omega$$

where x is a vector of inputs and ω is the vector of weights (or a matrix for a multi-class outcome).

- ▶ The predictive performance of perceptrons improved substantially by stacking them into multiple layers:



- ▶ Input variables are connected to multiple neurons in the hidden layer(s), which in turn are connected to output layer.
 - ▶ This is called a multi-layer perceptron or a feed forward neural network; with enough neurons, it can approximate any continuous function.
 - ▶ This is the “deep” in deep learning!

DNN Notation

- ▶ A DNN with a single hidden layer can be written as

$$y = \alpha_2 + g(\alpha_1 + x'\omega_1)'\omega_2$$

- ▶ α_1 and ω_1 , the intercept and coefficients in the input layer
 - ▶ ω_1 is a $d_0 \times d_1$ matrix, where d_0 is the dimension of the input and d_1 is the number of neurons in the hidden layer.
- ▶ $g(\cdot)$, the non-linear activation function.
 - ▶ without this, the DNN could only represent linear transformations of the input.
- ▶ α_2 and ω_2 , the intercept and coefficients in the hidden layer.
 - ▶ ω_2 is a $d_1 \times d_2$ matrix, where d_2 is the dimensionality of the output.



DNN Notation: two hidden layers

- ▶ Similarly, with two hidden layers we have

$$y = \alpha_3 + g_2(\alpha_2 + g_1(\alpha_1 + x'\omega_1)'\omega_2)'\omega_3$$


- ▶ $g_1(\cdot)$ and $g_2(\cdot)$, activation functions for the first and second layers.
- ▶ α_3 and ω_3 , intercepts and coefficients for the second hidden layer.

Constructing the Last Layer



- ▶ MLPs will output a probability distribution across output classes.
 - ▶ can also output a real number, which would make a regression model.

Modern MLPs: New activation functions



- ▶ logistic function: $\sigma(z) = \frac{1}{1+\exp(-z)}$
- ▶ hyperbolic tangent function: $\tanh(z) = 2\sigma(2z) - 1$
 - ▶ ranges between -1 and 1 (rather than between 0 and 1, as the case with the logistic)
 - ▶ centered on zero, can speed up convergence
- ▶ ReLU (rectified linear unit) function: $\max\{0, z\}$,
 - ▶ deceptively simple, fast to compute, and very effective in practice
 - ▶ gradient does not saturate to zero for large values (but is flat below zero)

Google Developers Advice: MLP baseline for Text Classification

1. Calculate the number of samples/number of words per sample ratio.
2. If this ratio is less than 1500, tokenize the text as n-grams and use a simple multi-layer perceptron (MLP) model to classify them.
 - ▶ In the case of N-grams models, Google testers found that MLPs tended to out-perform logistic regression and gradient boosting machines.

Python Implementation

- ▶ See the Jupyter notebook for Keras examples.
 - ▶ has not been updated to Keras 2.0 yet.
- ▶ “Dense” layer is the DNN baseline – means that all neurons are connected.
- ▶ Output layer:
 - ▶ for binary classification, use `activation='sigmoid'`
 - ▶ for regression, do not use an activation function
 - ▶ for multi-class classification, use `activation=softmax'`

Loss function and metrics

- ▶ Loss function:
 - ▶ for binary classification, use `binary_crossentropy`
 - ▶ for regression, use `mean_squared_error`
 - ▶ for multi-class classification, use `sparse_categorical_crossentropy`
- ▶ Metrics:
 - ▶ for classification, can use accuracy
 - ▶ for regression, can define a custom metric (see accompanying code)

Tuning NN Hyperparameters

- ▶ Number of hidden layers:
 - ▶ having a single hidden layer will generally give decent results.
 - ▶ more layers with fewer neurons can recover hierarchical relations and complex functions
 - ▶ for text classification, try one or two hidden layers as a baseline.
- ▶ Number of neurons:
 - ▶ a common practice is to set neuron counts like a funnel, with fewer and fewer neurons at each level
 - ▶ or just pick 150 neurons per layer
 - ▶ overall, better to have too many neurons, and use regularization
- ▶ Activation functions:
 - ▶ ReLU works well for hidden layers
 - ▶ softmax is good for the output layer in classification tasks

Xavier and He Initialization

Activation function	Uniform distribution $[-r, r]$	Normal distribution
Logistic	$r = \sqrt{\frac{6}{n_{\text{inputs}} + n_{\text{outputs}}}}$	$\sigma = \sqrt{\frac{2}{n_{\text{inputs}} + n_{\text{outputs}}}}$
Hyperbolic tangent	$r = 4\sqrt{\frac{6}{n_{\text{inputs}} + n_{\text{outputs}}}}$	$\sigma = 4\sqrt{\frac{2}{n_{\text{inputs}} + n_{\text{outputs}}}}$
ReLU (and its variants)	$r = \sqrt{2}\sqrt{\frac{6}{n_{\text{inputs}} + n_{\text{outputs}}}}$	$\sigma = \sqrt{2}\sqrt{\frac{2}{n_{\text{inputs}} + n_{\text{outputs}}}}$

- Connection weights should be initialized randomly according to a uniform distribution or normal distribution, as indicated in the table (see Geron Chapter 11).

Other Activation Functions


- ▶ Leaky ReLU

$$\max(\alpha z, z)$$

where α is set to a small number, such as .01, or learned in training.

- ▶ Exponential linear unit

$$\text{ELU}(z) = \begin{cases} \alpha(\exp(z) - 1) & z < 0 \\ z & z \geq 0 \end{cases}$$



- ▶ In general, ELU has had the best performance so far, but it is slower than ReLU.

Batch normalization

- ▶ Another trick to speed up training:
 - ▶ in between layers, zero-center and normalize the inputs to variance one.
 - ▶ normally done before a non-linear activation function

Regularization for Sparse Models

- ▶ As with linear models, neural network parameters can be regularized with an L1 and/or L2 penalty to push weak neurons to zero and produce a sparse model.

Dropout

- ▶ An elegant regularization technique:
 - ▶ at every training step, every neuron has some probability (typically 0.5) of being temporarily dropped out, so that it will be ignored at this step.
 - ▶ after training, neurons don't get dropped any more.
- ▶ Neurons trained with dropout:
 - ▶ cannot co-adapt with neighboring neurons and must be independently useful.
 - ▶ cannot rely excessively on just a few input neurons; they have to pay attention to all input neurons.
 - ▶ makes the model less sensitive to slight changes in the inputs.
- ▶ If a model is over-fitting, increase dropout. Dropout can be higher for large layers and lower for small layers.

Optimizers

- ▶ Choice of optimization algorithm is the topic of active research, which has shown that it can have a big impact on model performance.
 - ▶ Until recently, a good starting choice would be Adam (adaptive moment estimation), which is fast and usually works well. For robustness, can also try SGD.
 - ▶ A recent paper says that AdaBound dominates Adam or SGD.

Early stopping

- ▶ A popular/efficient regularization method is to continually evaluate your model at regular intervals, and then to stop training when the test-set accuracy starts to decrease.

Practical Guidelines

Table 11-2. Default DNN configuration

Initialization	He initialization
Activation function	ELU
Normalization	Batch Normalization
Regularization	Dropout
Optimizer	Adam
Learning rate schedule	None

Source: Geron book.

Batch Training with Large Data

- ▶ If data sets don't fit in memory, one can load the data in batches from disk.
- ▶ can also continuously update a saved model.

Grid search for model choice

- ▶ The flexibility of DNNs is a blessing and a curse.
 - ▶ in general, one should make a complex model that allows regularization.
- ▶ But still, there are many choices to be made.
 - ▶ to choose the number of hidden layers, for example, one can use cross-validation grid search (as we did with standard scikit-learn models).